



An Architects Guide to Delivering Data Insights

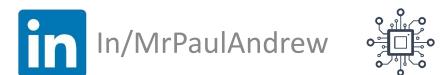
Using the Microsoft Azure Data Platform

Paul Andrew | Principal Consultant & Solution Architect



Microsoft®
Most Valuable
Professional









Thank you to our Global Sponsors





IDERA

Quest

Sentry One.

vmware[®]



Thank you to our local Partners





















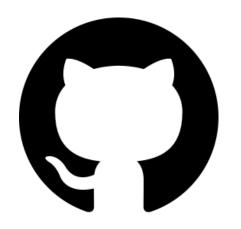












https://github.com/mrpaulandrew

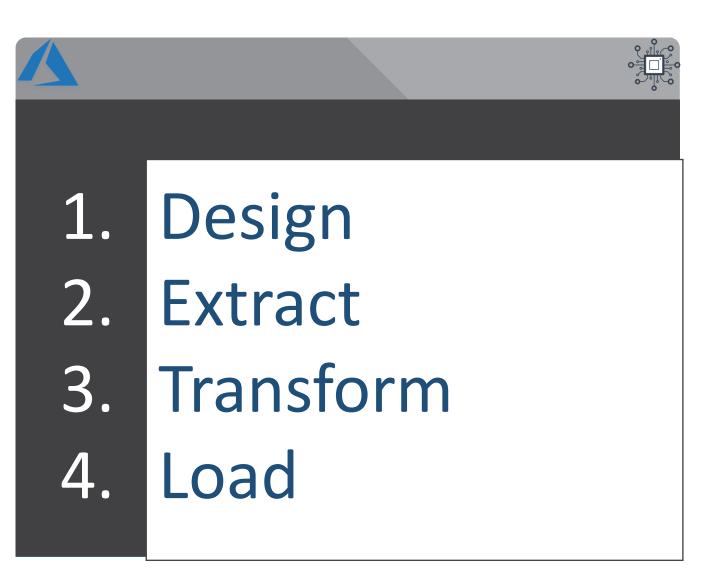
CommunityEvents

Demo code, content and slides from various community events.

C++

{Event/Location}-{Month}-{Year}

Agenda



Question:

What is the answer to life, the universe and everything?



Answer: It depends!





Question:

What is big data?

Answer:

It depends!



Answer:

Any data that you cannot process in the time that you have using the technology you have.

- Buck Woody







- Disaster recovery
- Transaction level restart ability

Resilience

Rapid Delivery

- Metadata driven
- Continuous deployments

(Re)Usability

- Generic code libraries
- **Dataset contracts**



Performance

- Complex partitioning
- Large compute clusters

Scalability

- Microservices
- Small discrete batches

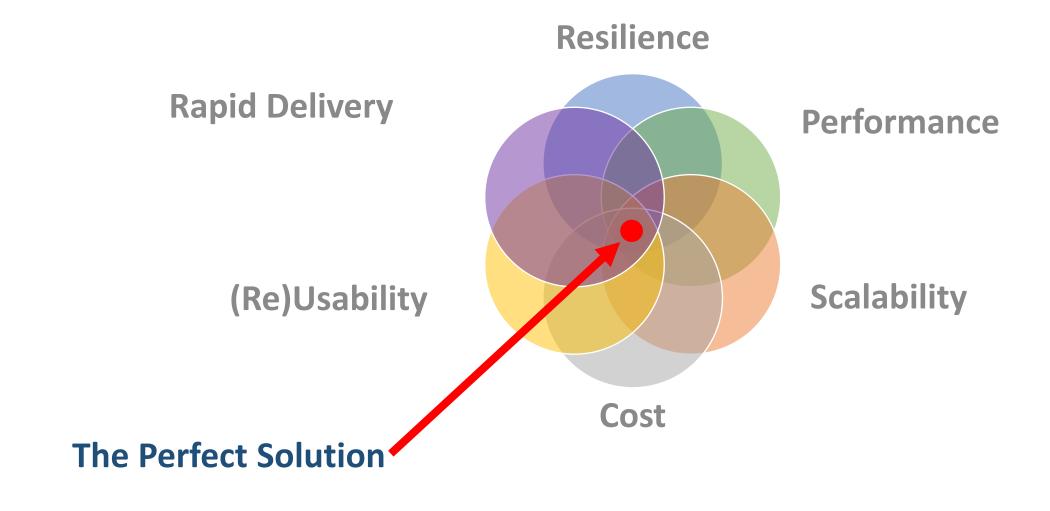
Cost

- Minimum resources used
- Careful resource management

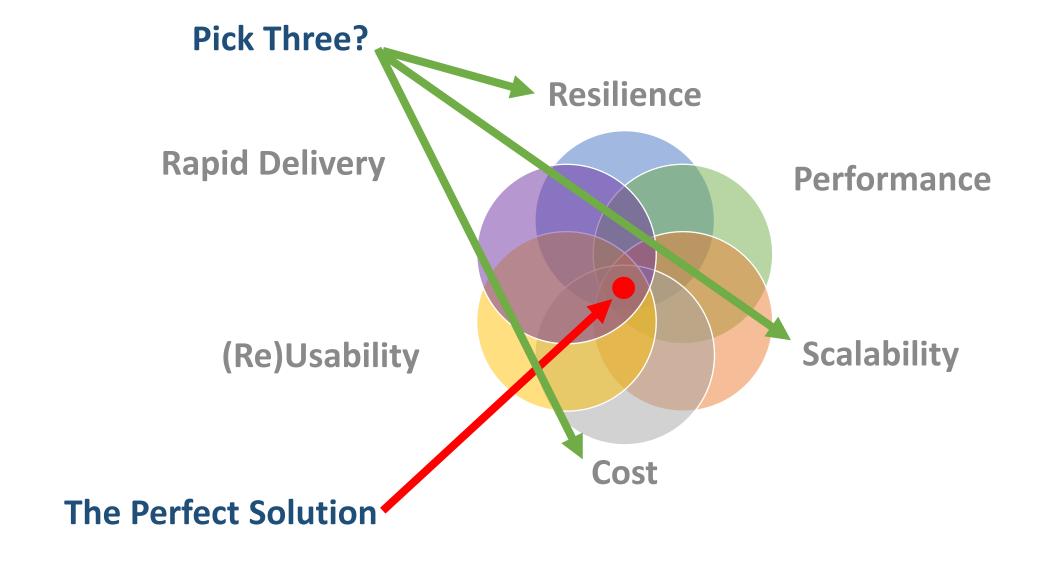






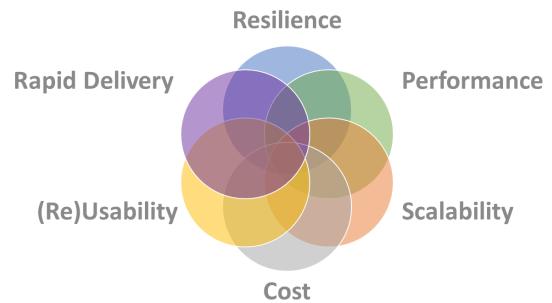






Agenda







Data Extraction & Ingestion



Data Structure



Data Source











Push or Pull





Batch or Stream









Public or **Private Transfer**









Data Sensitivity









Data Volume

























Push or Pull





Batch or Stream

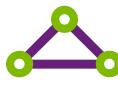






Public or **Private Transfer**









Data Sensitivity







Data Volume

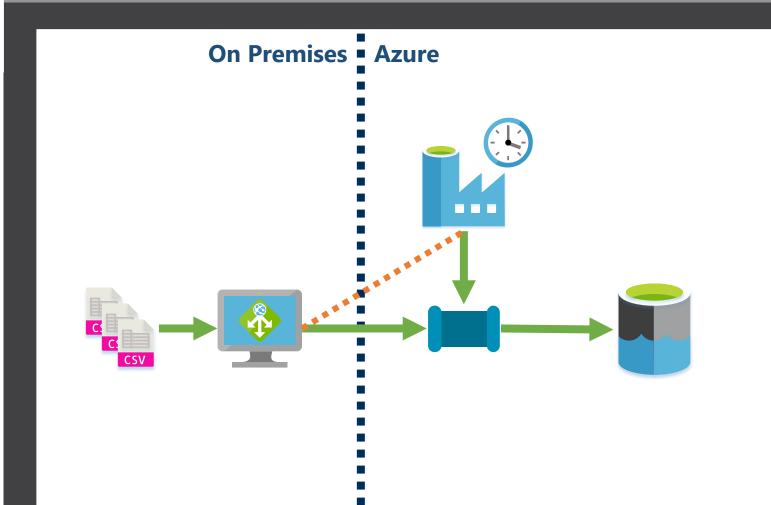












Requirements:

- Flat files
- From local storage
- Pulled from source
- Batch load
- Public connections
- No PII data
- Small data volumes









































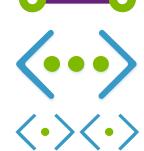










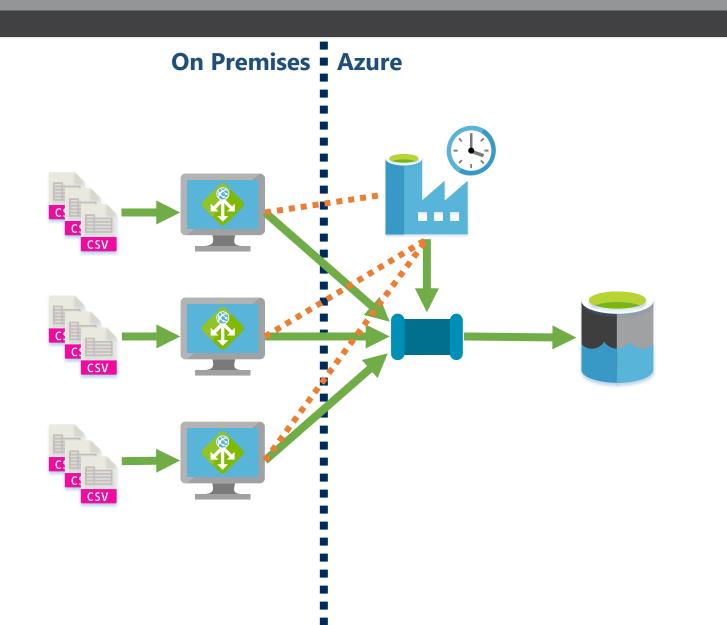












Requirements:

- Flat files
- From local storage
- Pulled from source
- Batch load
- Public connections
- No PII data
- <u>Large</u> data volumes

























Push or Pull





Batch or Stream



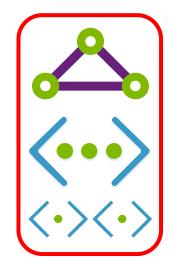




Public or **Private Transfer**







Data Sensitivity

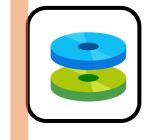






Data Volume

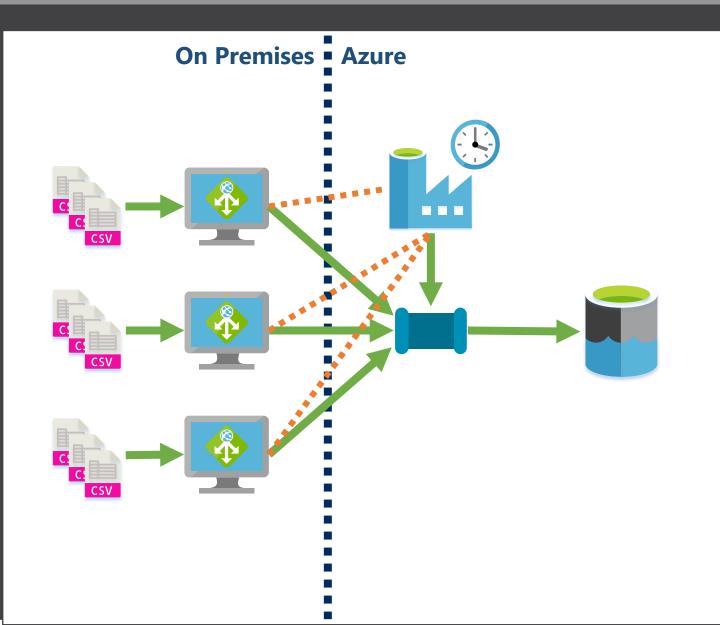










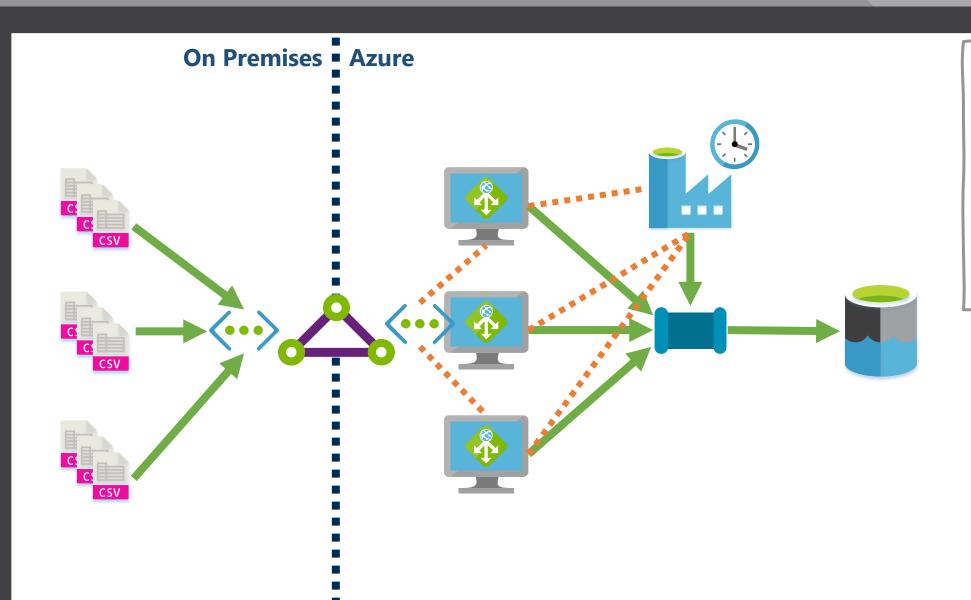


Requirements:

- Flat files
- From local storage
- Pulled from source
- Batch load
- Private connections
- No PII data
- Large data volumes







Requirements:

- Flat files
- From local storage
- Pulled from source
- Batch load
- Private connections
- No PII data
- Large data volumes





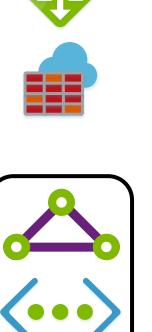






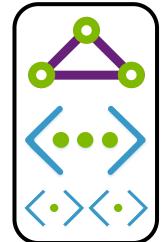


Batch or Stream



Public or **Private Transfer**





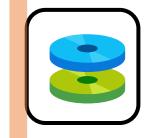
Data Sensitivity





Data Volume















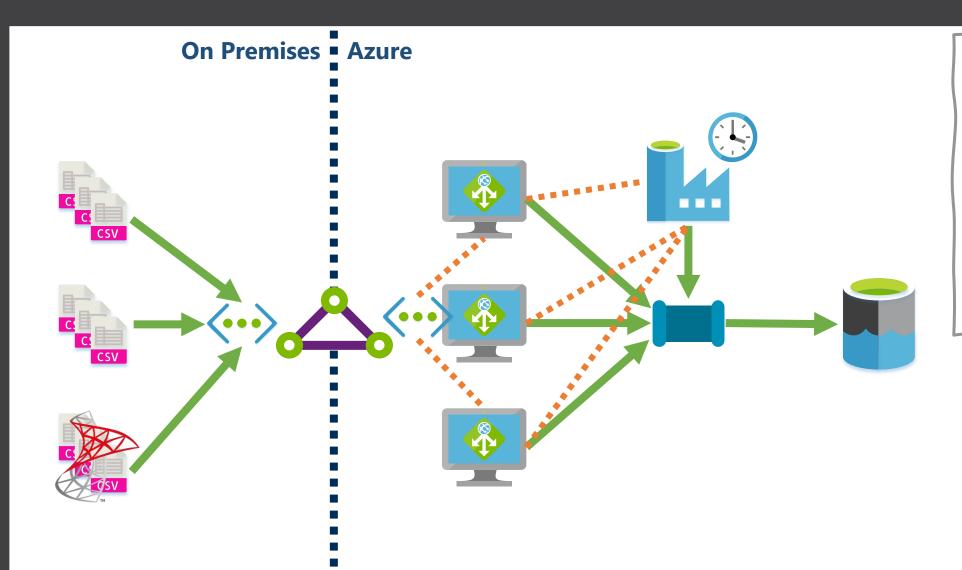












Requirements:

- Flat files
- From local storage & database tables
- Pulled from source
- Batch load
- Private connections
- No PII data
- Large data volumes









Batch or Stream



























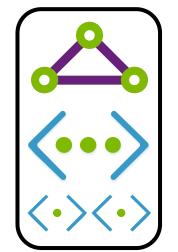








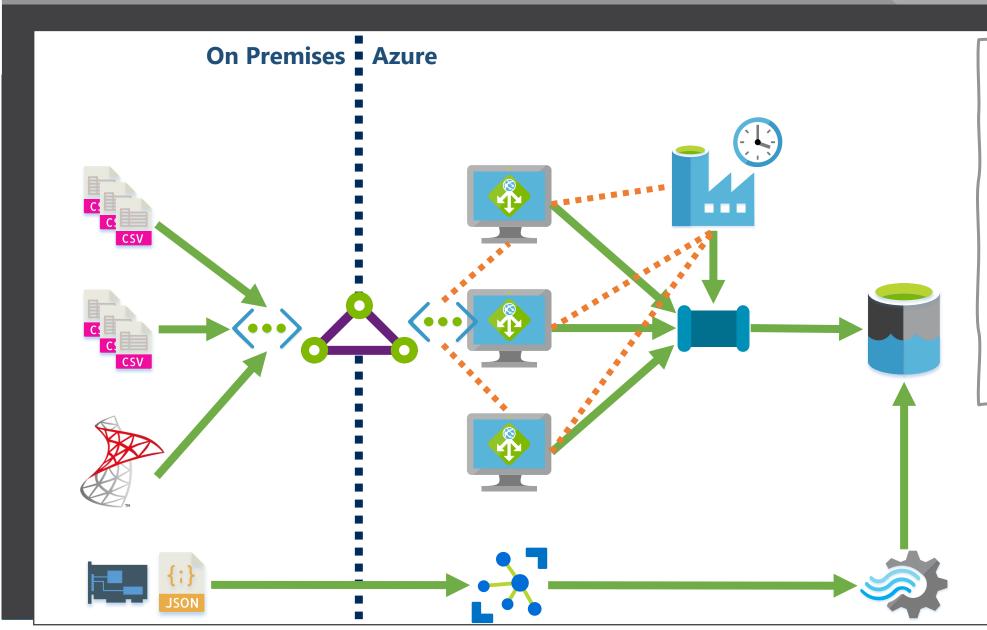










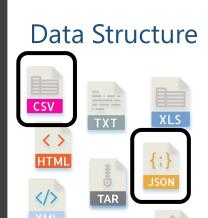


Requirements:

- Flat files & JSON
- From local storage & database tables
- Pulled from source& pushed
- Batch load & streamed
- Private connections
- No PII data
- Large data volumes













Batch or Stream

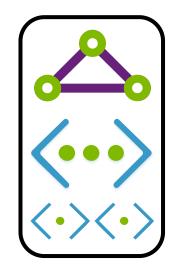




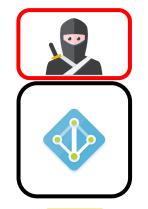
Public or **Private Transfer**







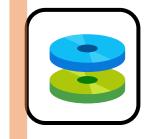
Data Sensitivity





Data Volume





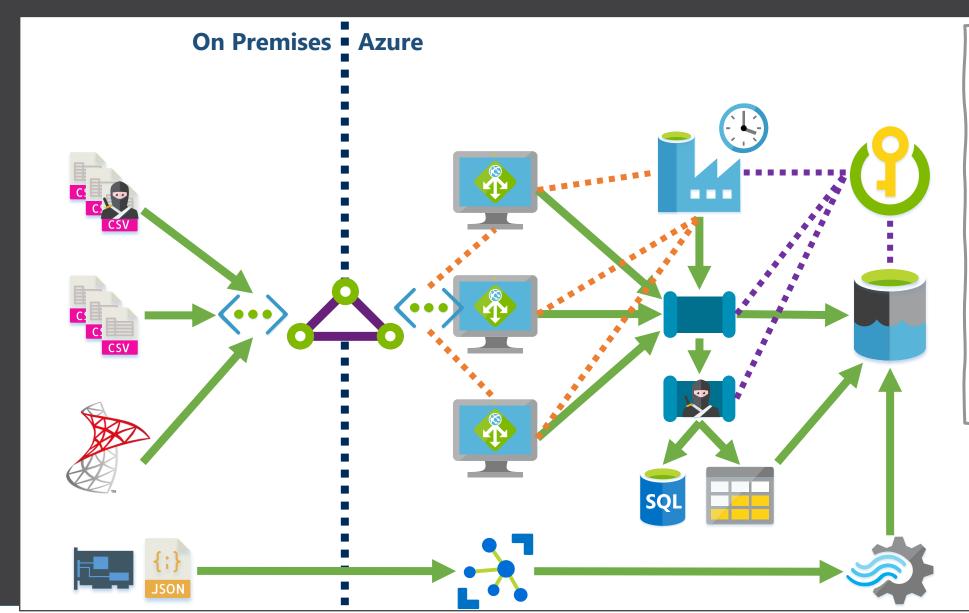












Requirements:

- Flat files & JSON
- From local storage & database tables
- Pulled from source& pushed
- Batch load & streamed
- Private connections
- Both PII data & none
 PII data
- Large data volumes



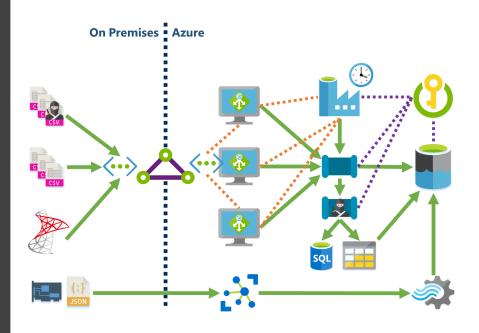
Overall Architecture



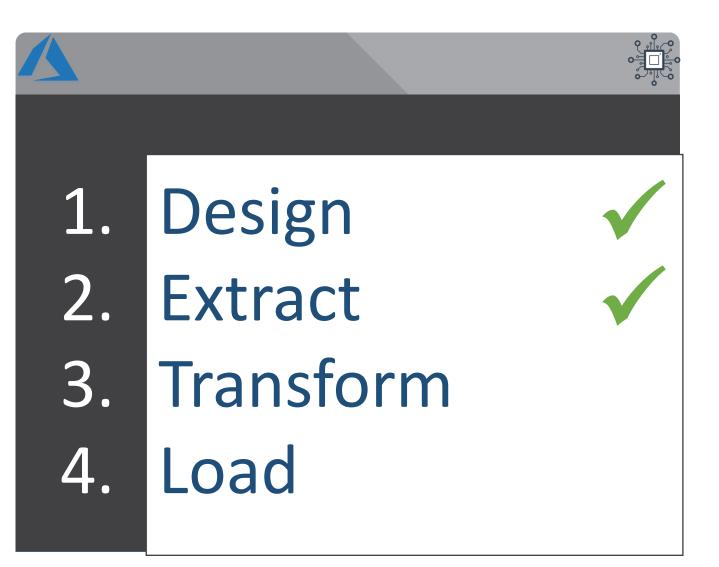
Extract

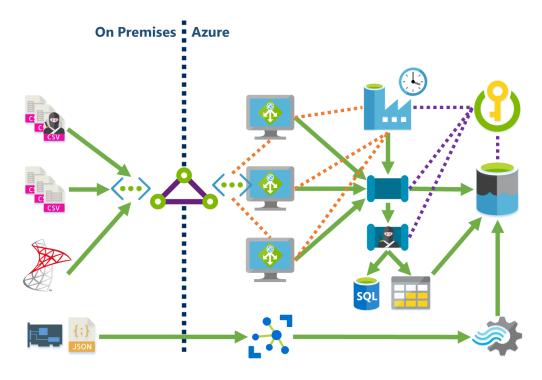
Transform

Load

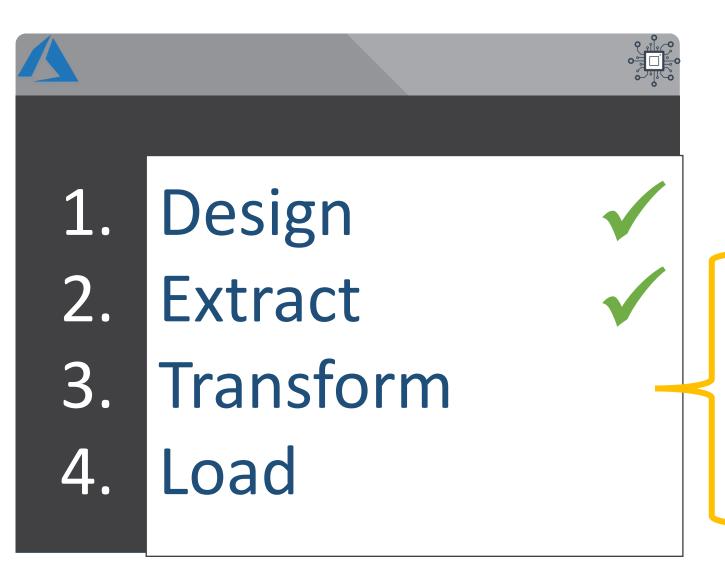


Agenda





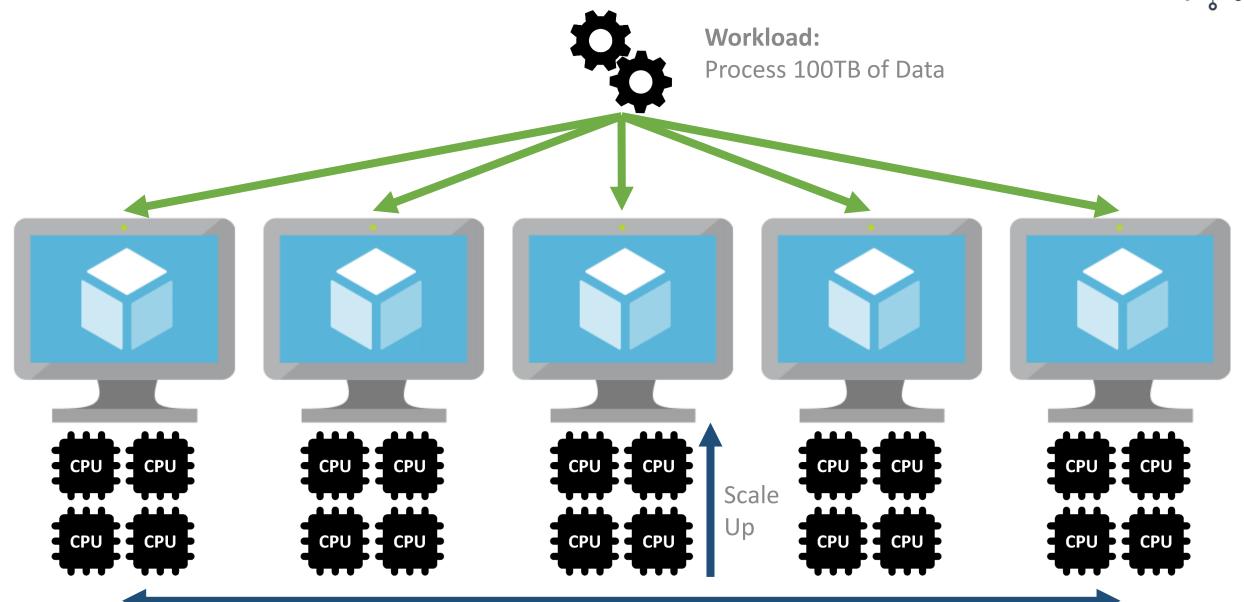
Agenda



Compute
Storage, Structure
& Data Format

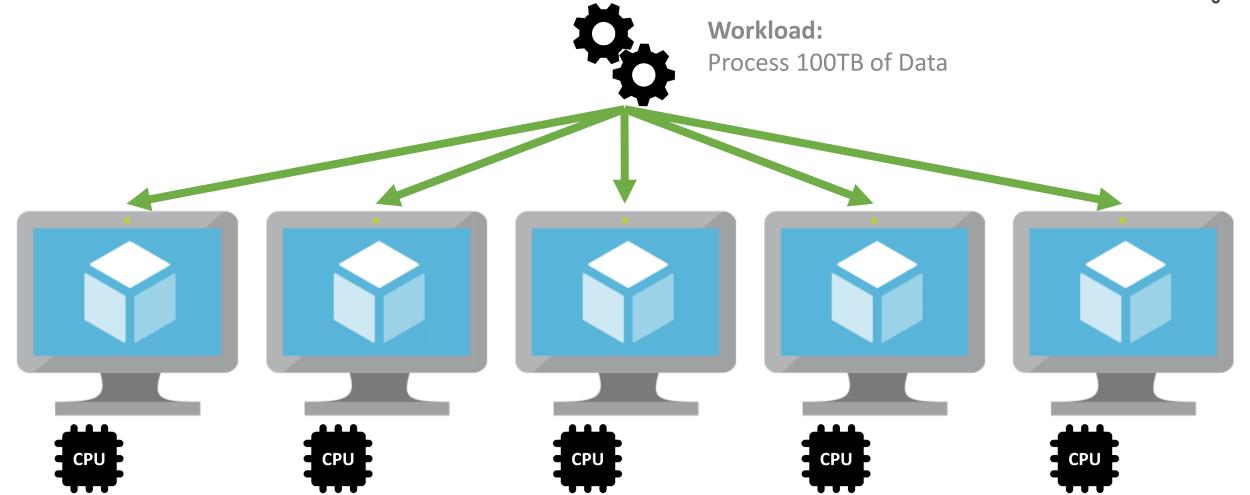
Scaling Up and/or Scaling Out





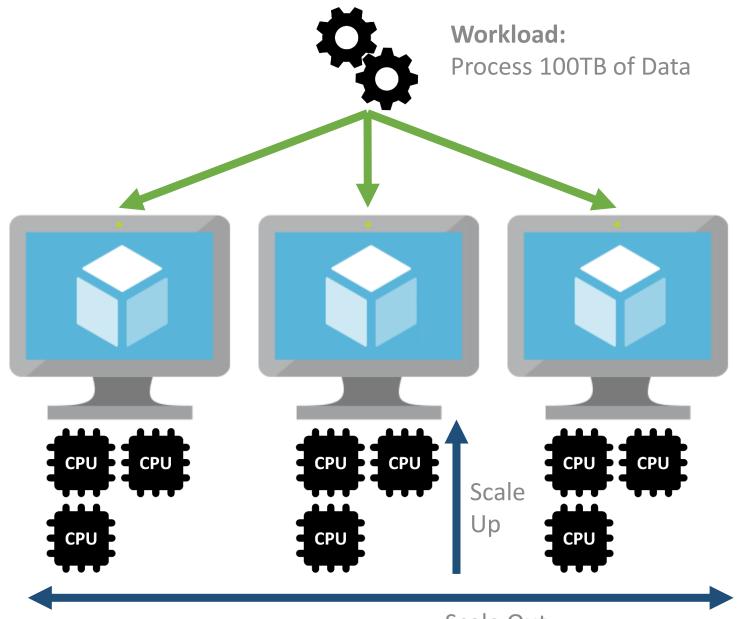
Scaling Up and/or Scaling Out





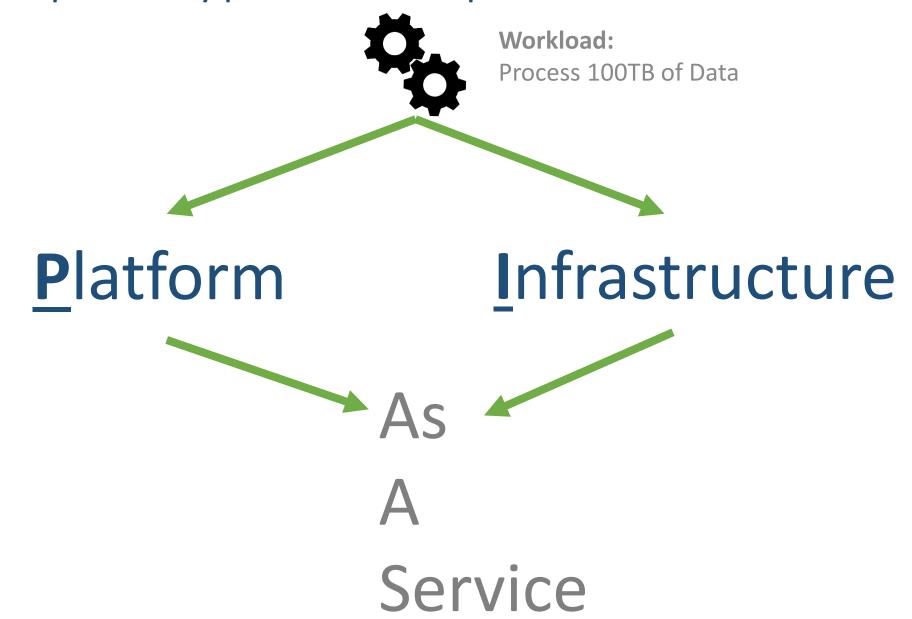
Scaling Up and/or Scaling Out





What Compute Type of Compute?





What Compute Type of Compute?





Infrastructure

As

A

Service



Data Transformation – Compute



Database **Engine**

Integration Services

Analysis Services



Reporting Services

Data Quality Services

Master Data Services



Data Transformation – Compute



2019Big Data Cluster



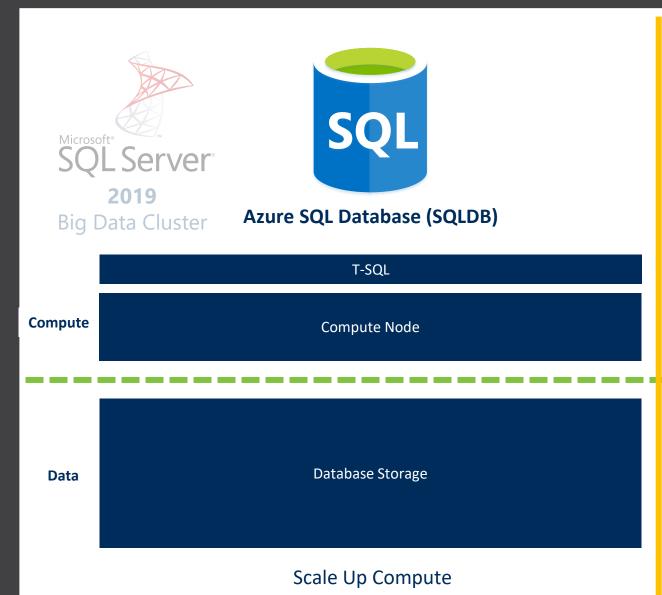










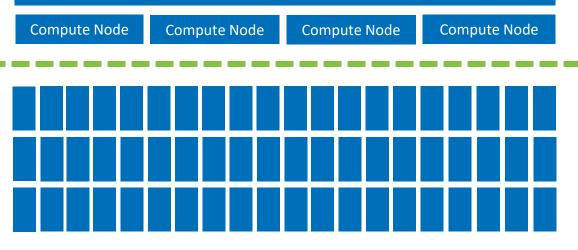




Azure SQL Data Warehouse (SQLDW)

D-SQL

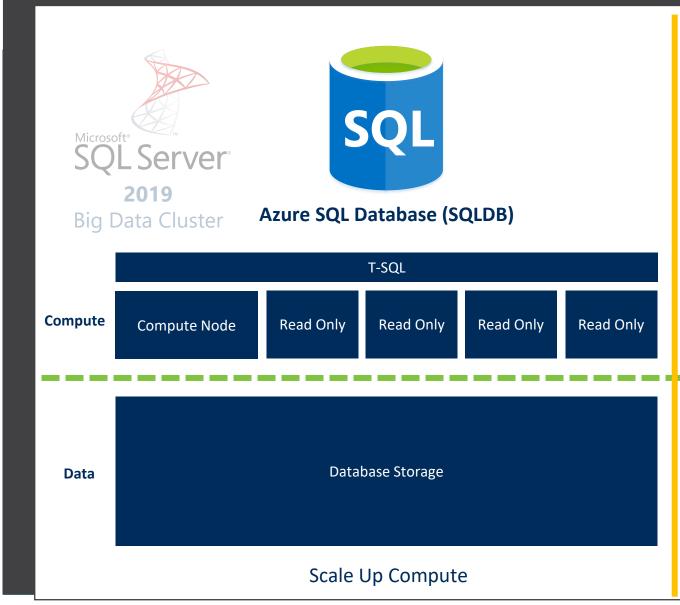
Control Node



Scale Out Compute & Distributed Storage



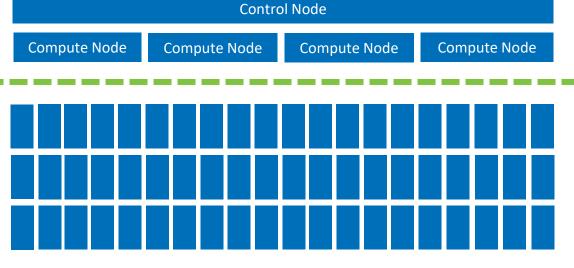






Azure SQL Data Warehouse (SQLDW)

D-SQL

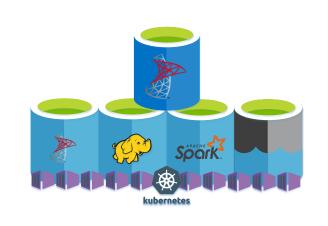


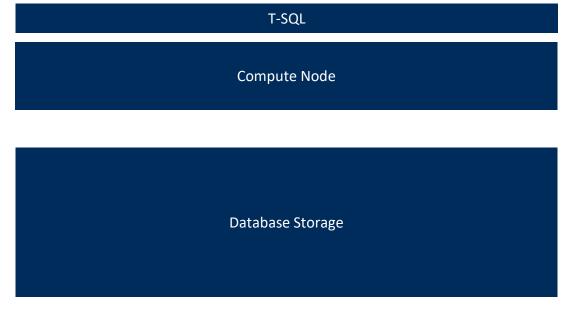
Scale Out Compute & Distributed Storage

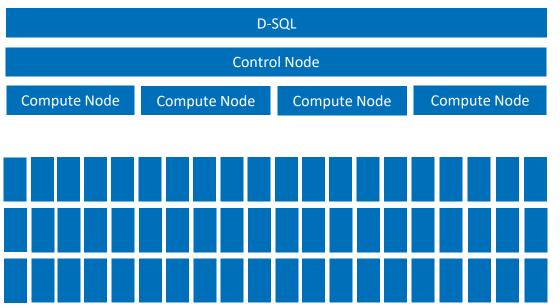






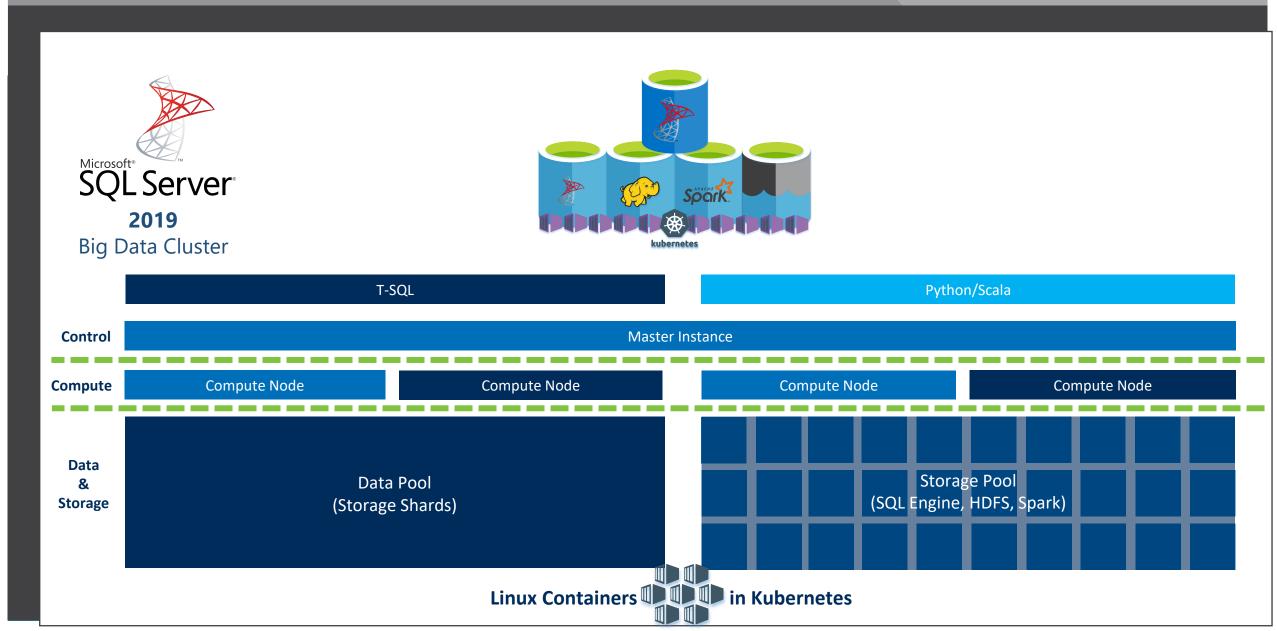








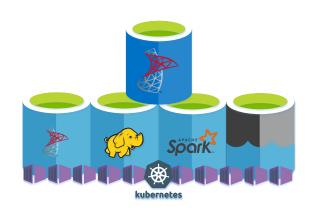








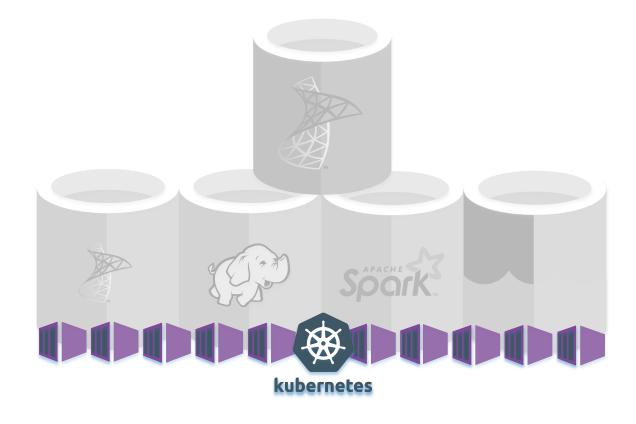










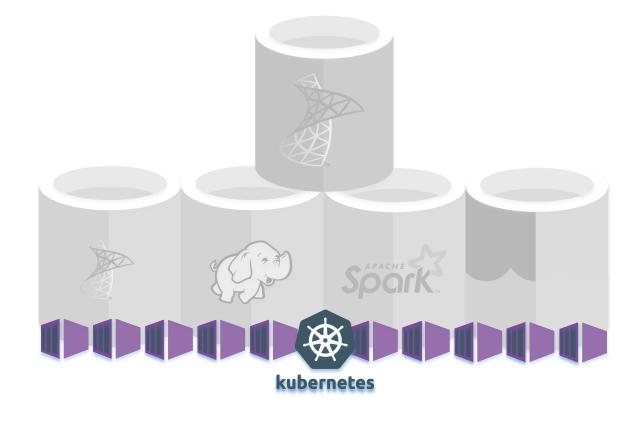












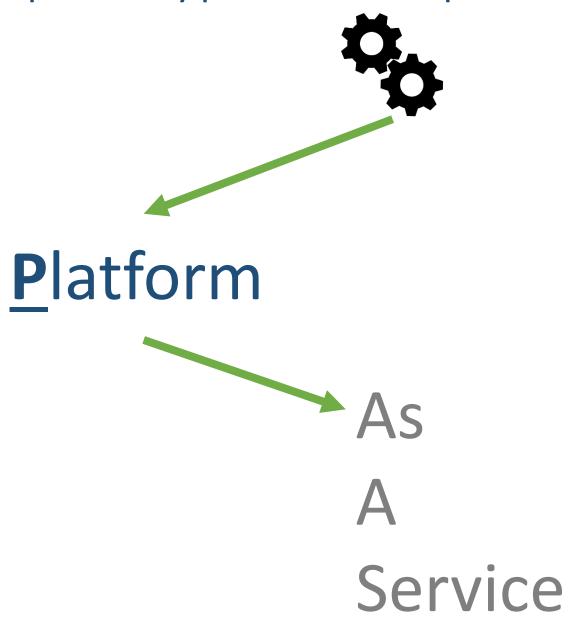






What Compute Type of Compute?









Data Lake Analytics

HDInsight

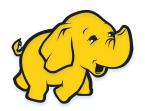
SQL Database SQL Data Warehouse (Synapse)

Databricks

Custom Activity
Batch Service

SSIS Packages















Automation



Cosmos

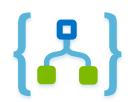




Power BI Data Flows



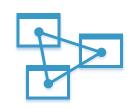
Logic Apps



Data Flows



Analysis Services





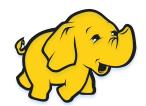


Data Lake Analytics



U-SQL to be deprecated?

HDInsight



Manual cluster configuration.

nt SQL Database



Only scales up. Hyperscale tier.

SQL Data Warehouse (Synapse)



Scales out. Is Synapse ready?

Databricks



Learn Scala. As open as Spark?

Custom Activity
Batch Service



Build your own! Scales up.



SSIS Packages

laaS to run. Scales up.

Automation



Short amounts of compute.

Cosmos



Transactional data. Slow with batch data.

Functions



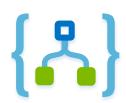
Build your own! Durable?

Power BI Data Flows



Not enterprise grade.

Logic Apps



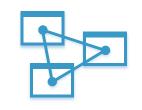
Event based actions & data movements.

Data Flows



Currently limited transformations.

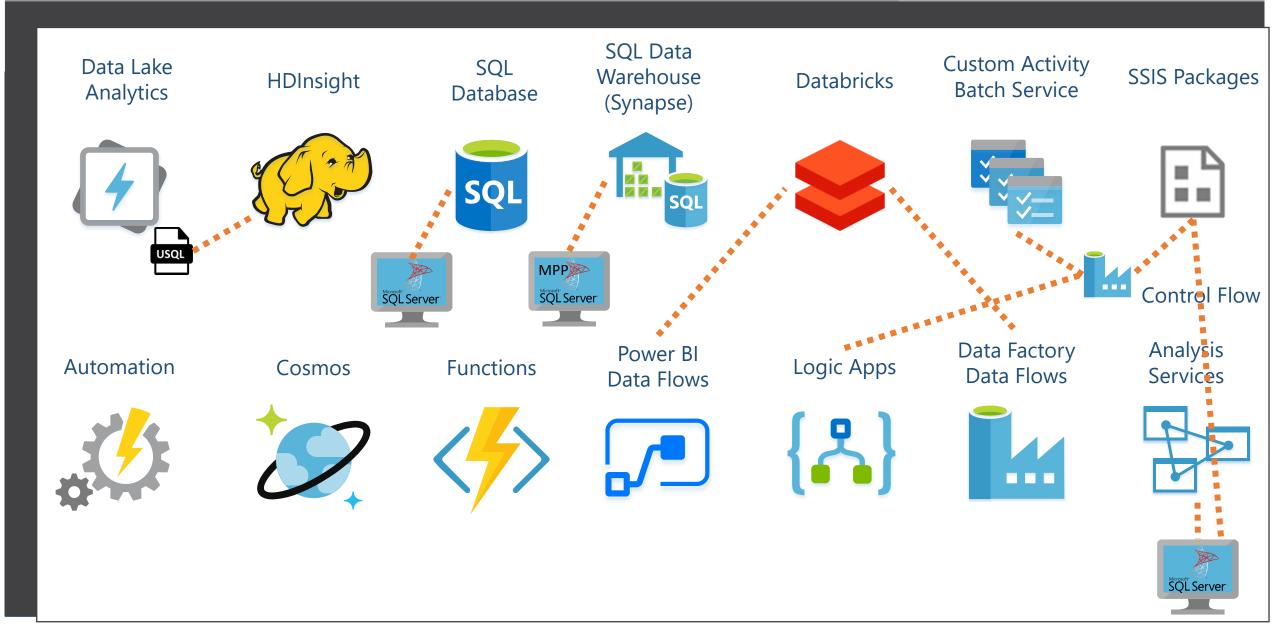
Analysis Services



Calculation optimised. Not for engineering.











Data Lake Analytics

HDInsight

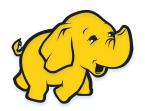
SQL Database SQL Data Warehouse (Synapse)

Databricks

Custom Activity
Batch Service

SSIS Packages















Automation



Cosmos

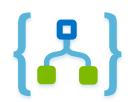




Power BI Data Flows



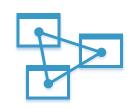
Logic Apps



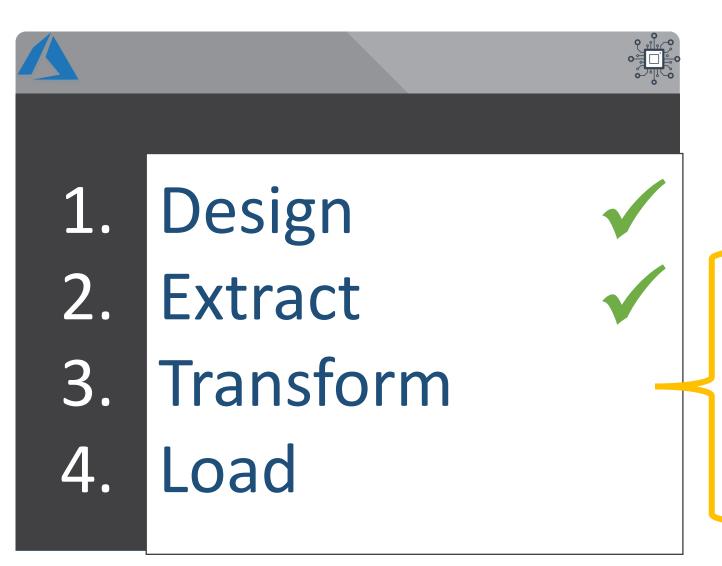
Data Flows



Analysis Services



Agenda



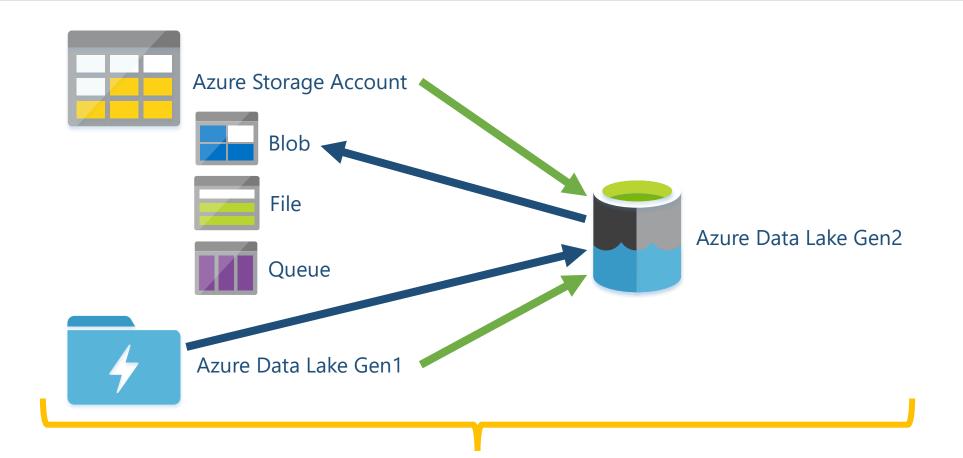
Compute

Storage, Structure

& Data Format







Hadoop Distributed File System (HDFS)





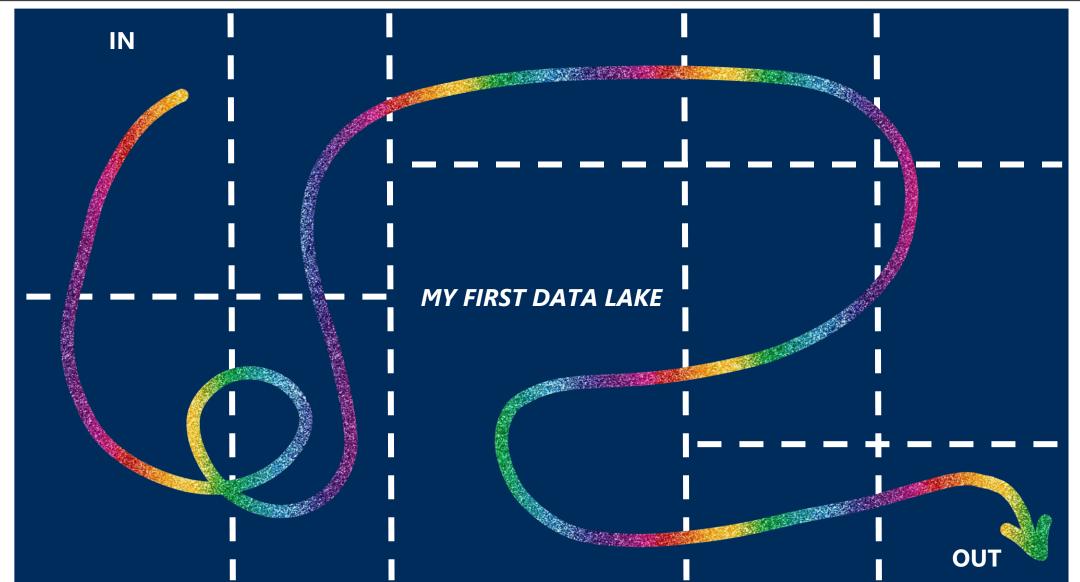






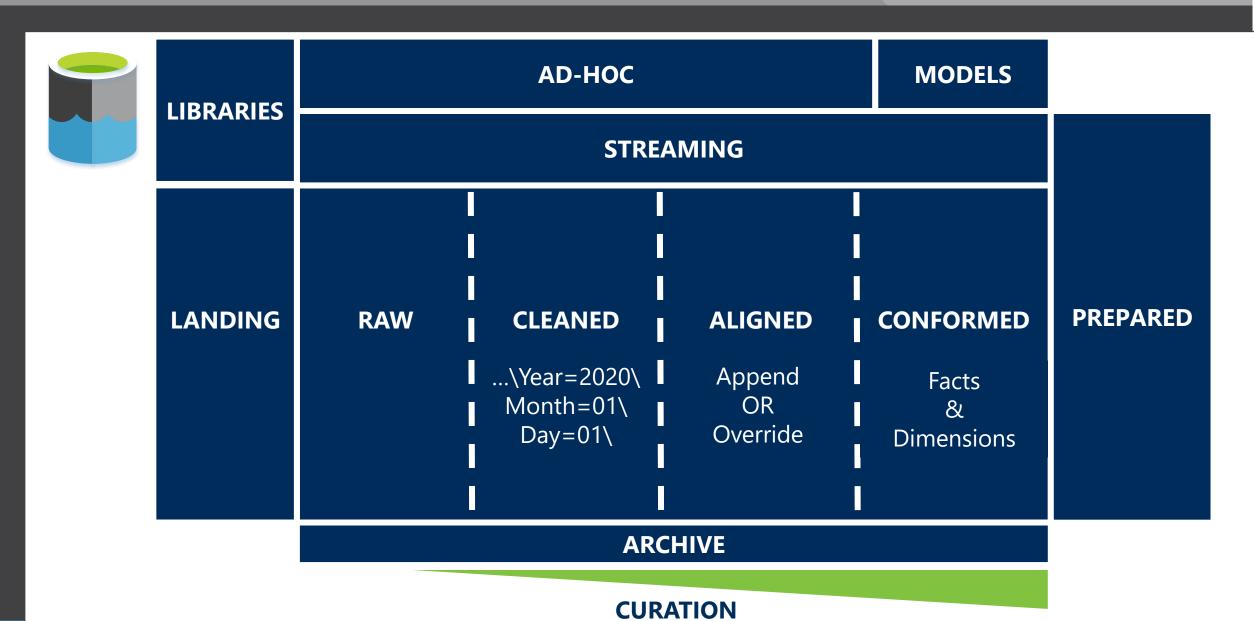




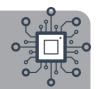


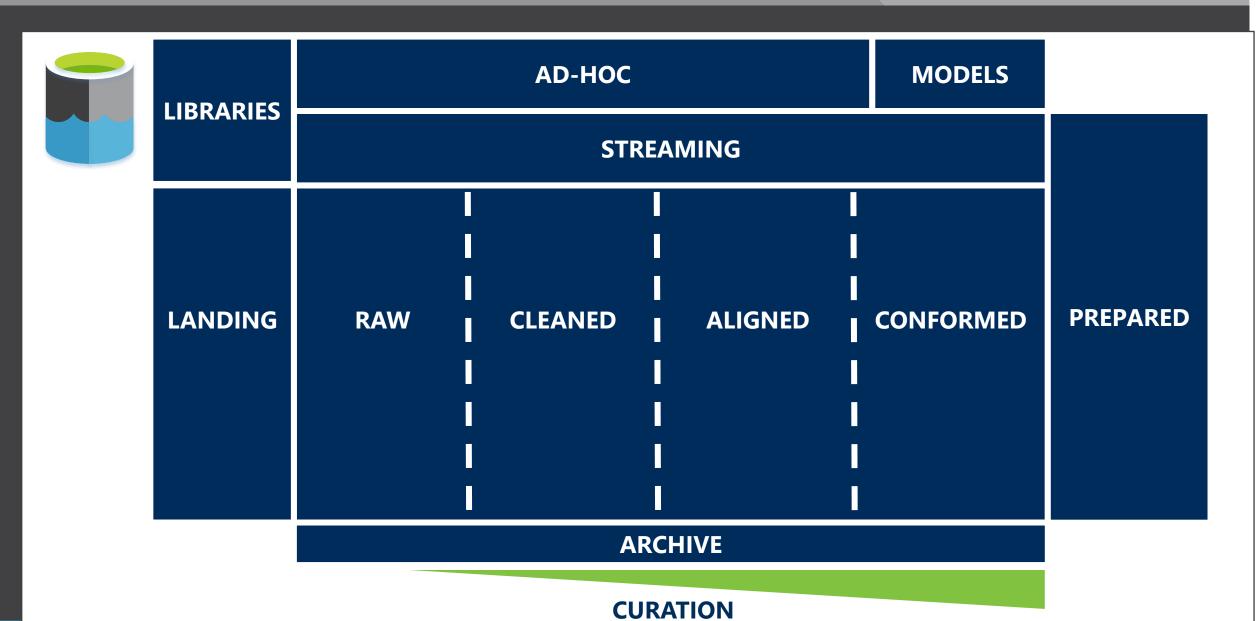






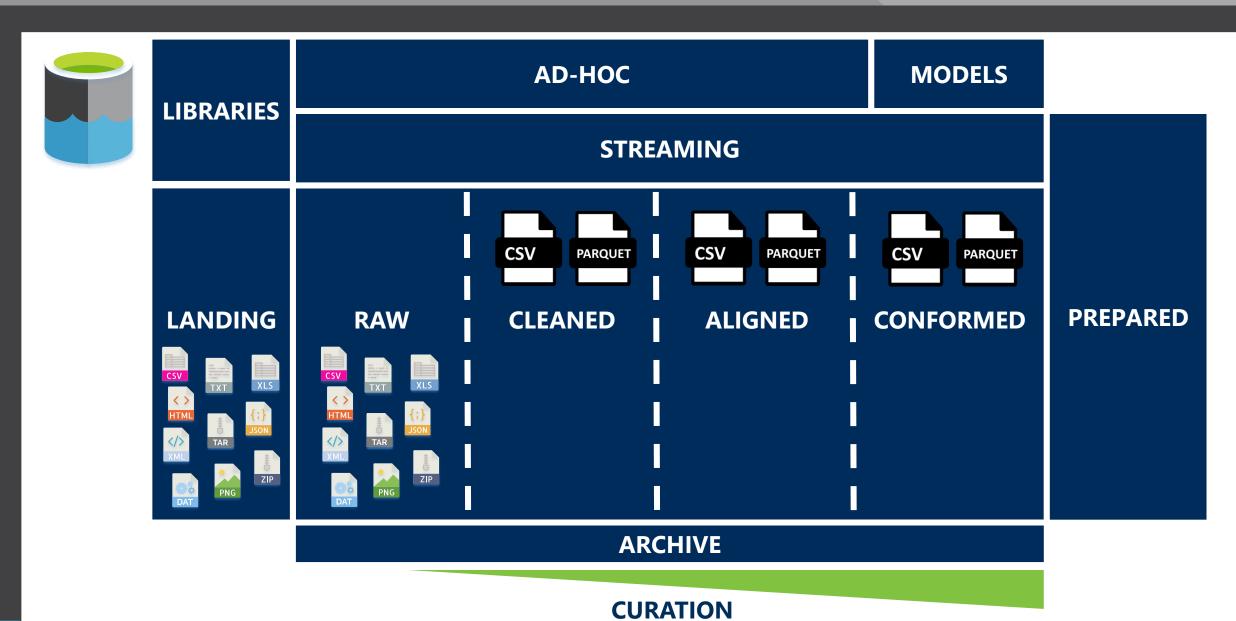






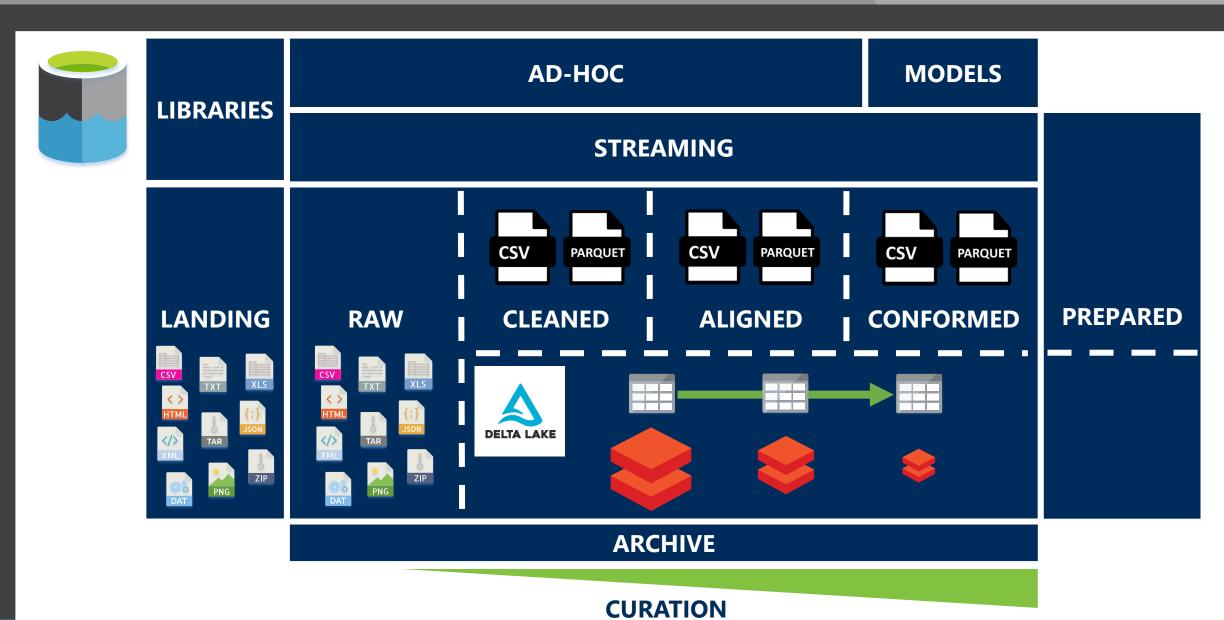




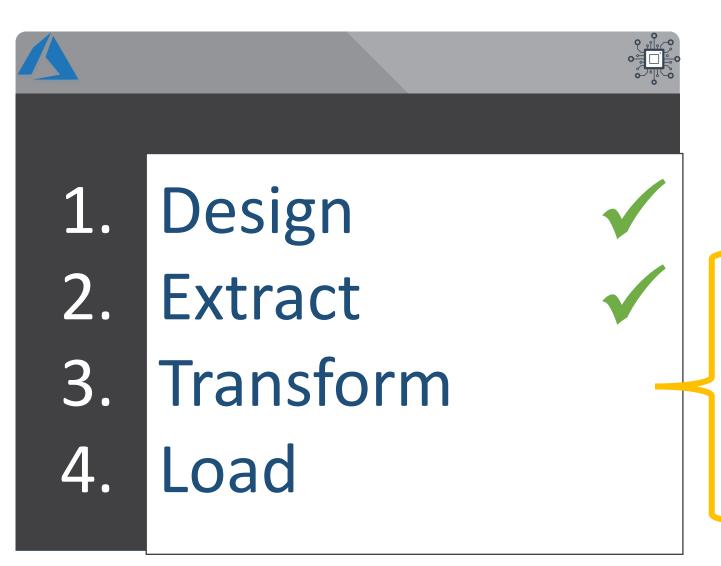








Agenda



Compute

Storage, Structure

& Data Format



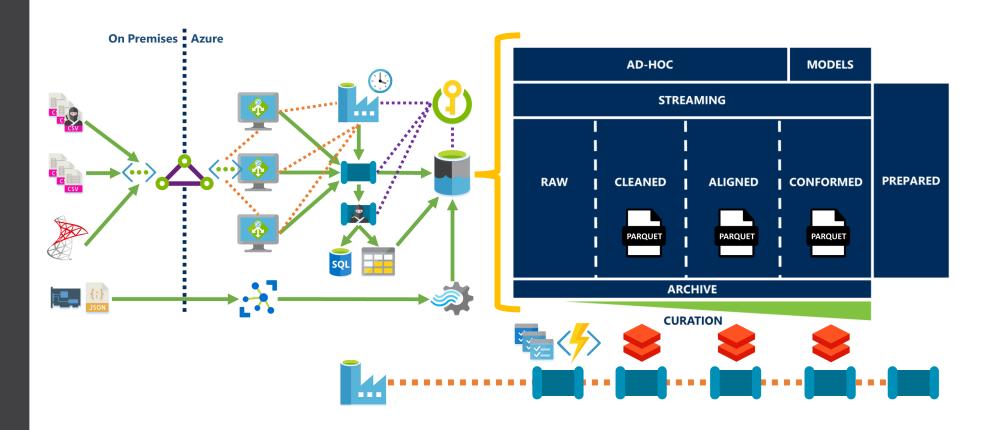
Overall Architecture



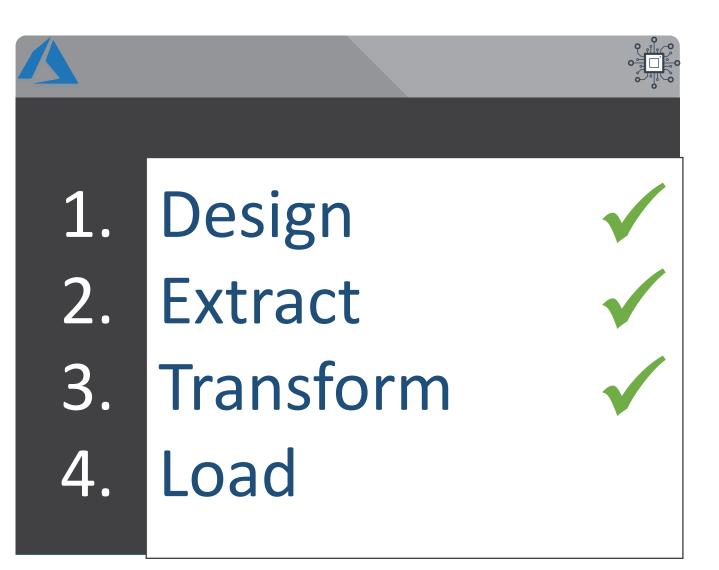
Extract

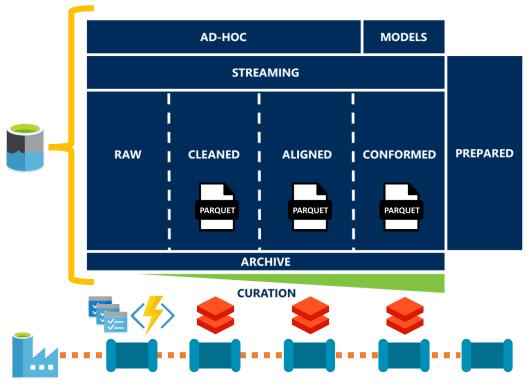
Transform

Load



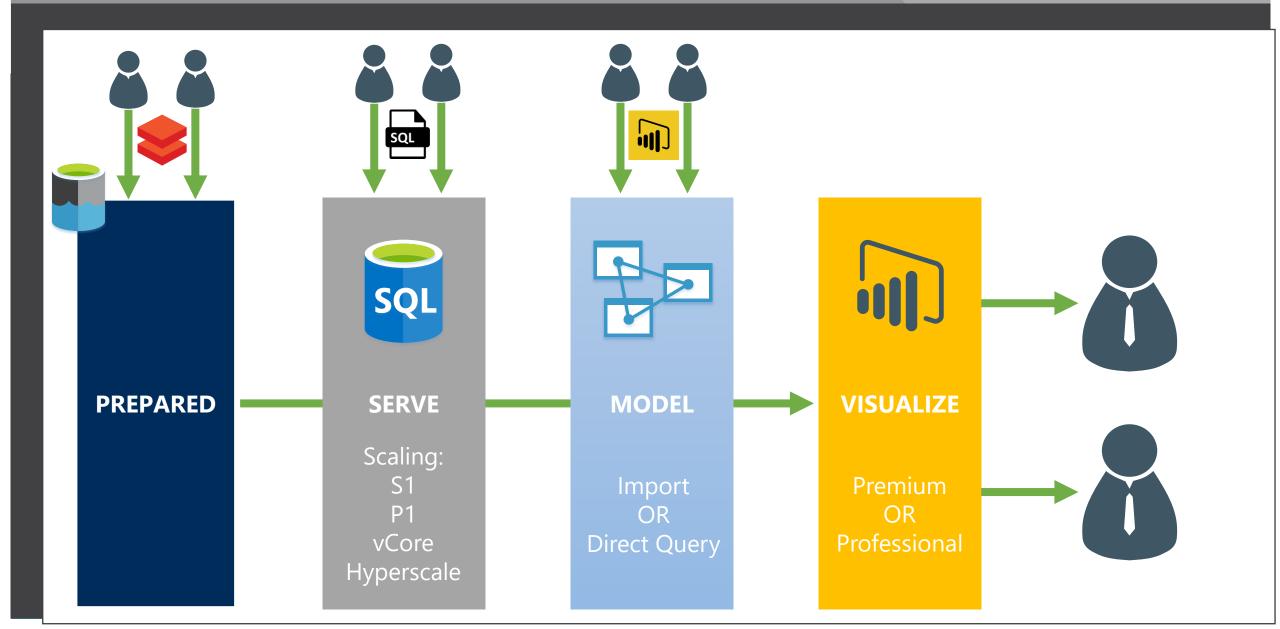
Agenda





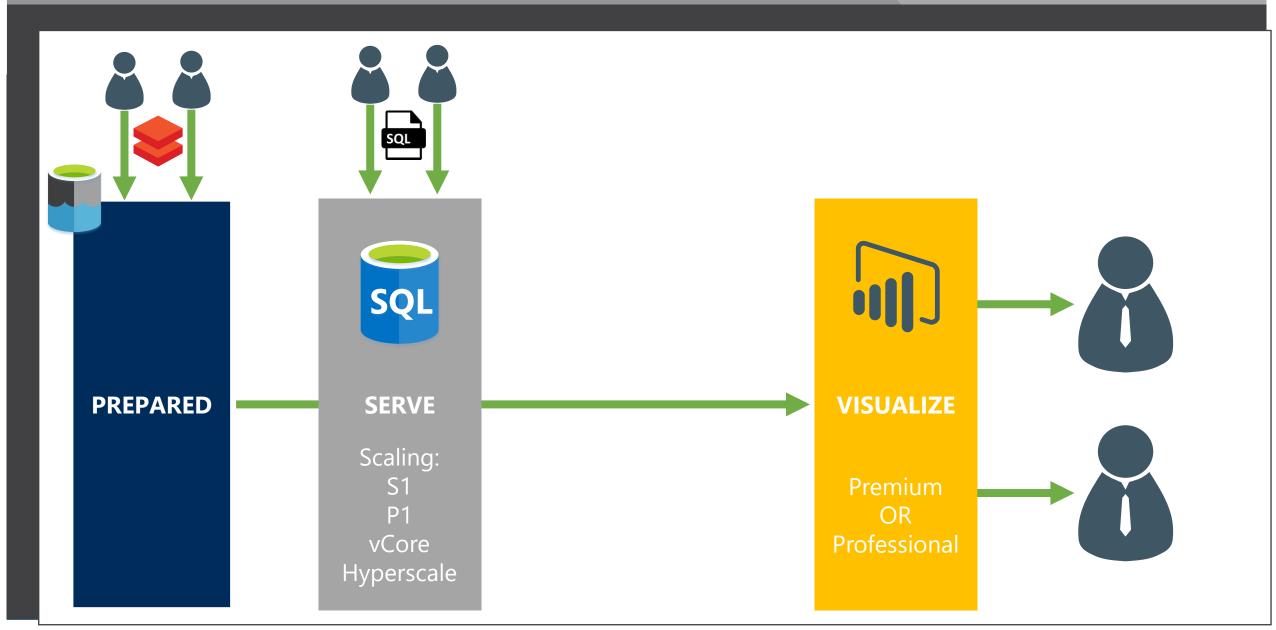






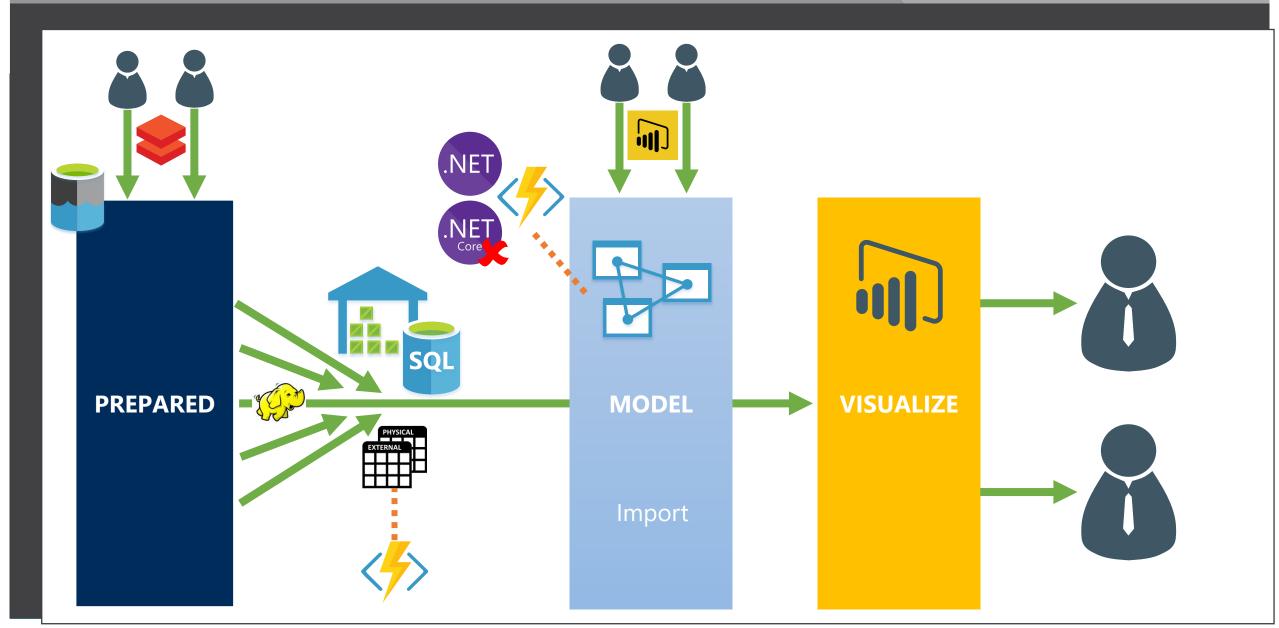






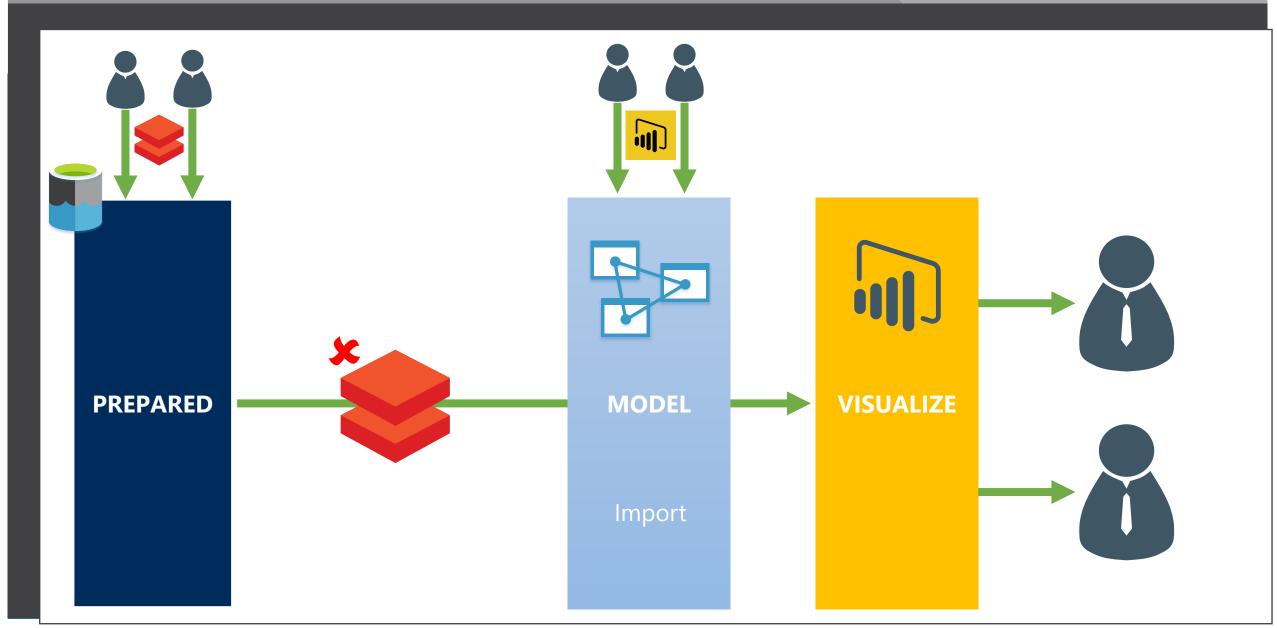






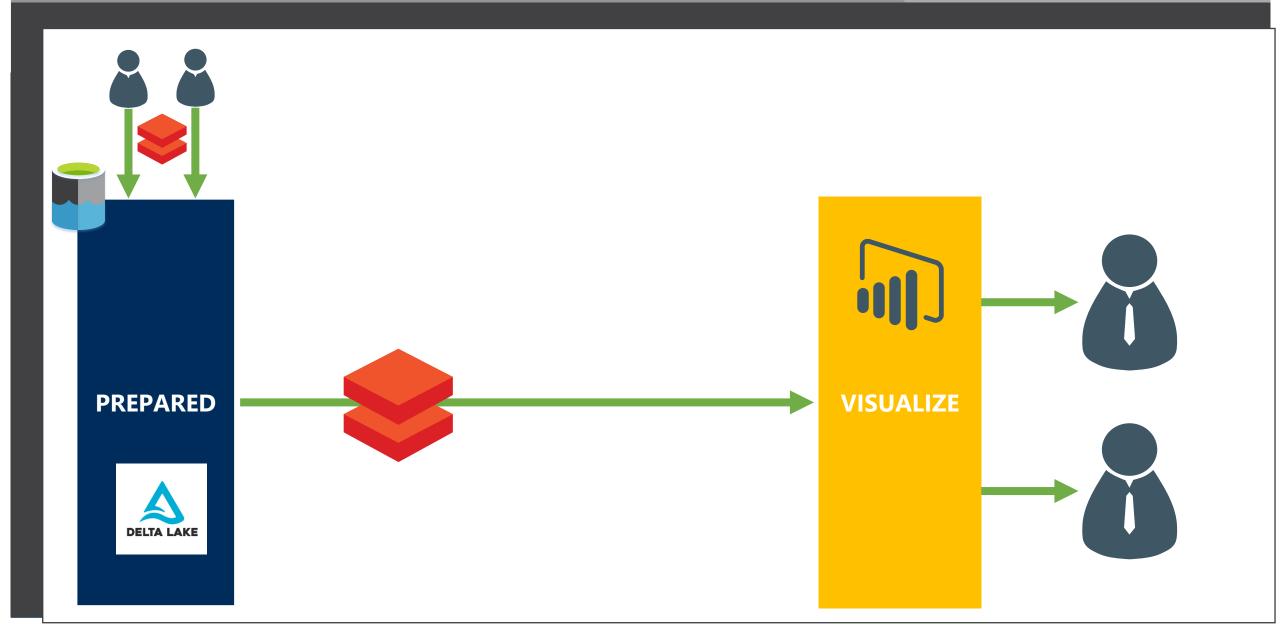














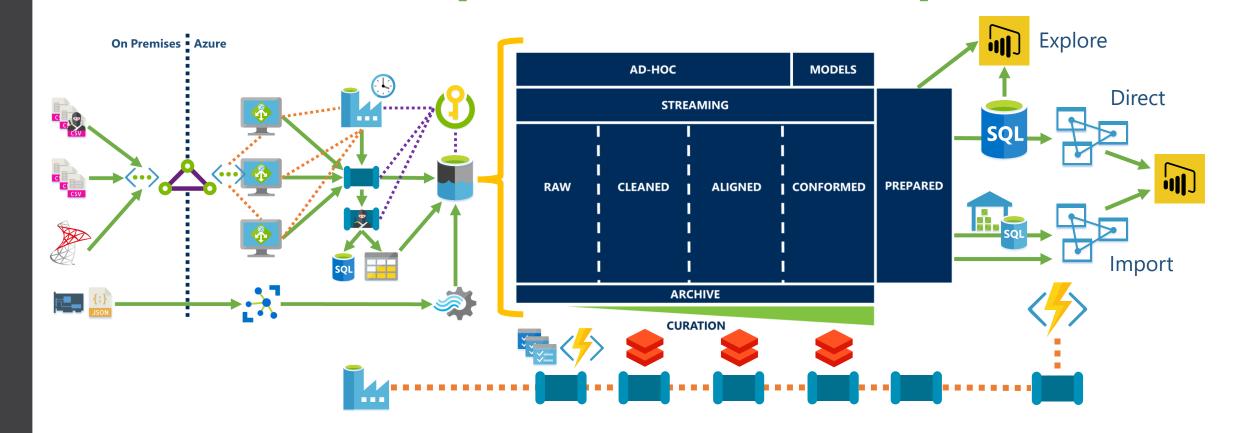
Overall Architecture



Extract

Transform

Load





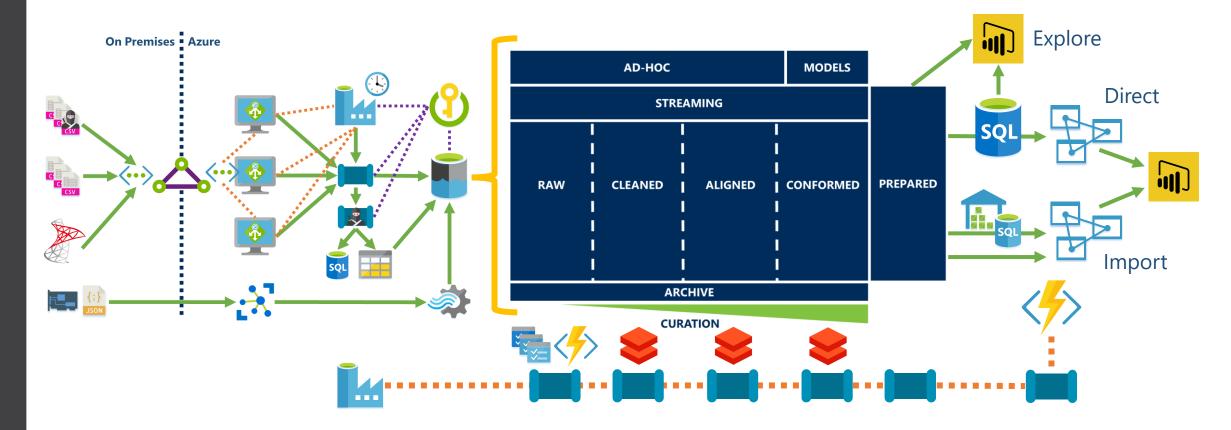
Overall Architecture



Extract

Transform

Load



Q: Should we build our data platform solution like this?...A: It depends!





Thank you for listening...



Paul Andrew Myp Microsoft® Most Valuable Professional Altius

mrpaulandrew.com Blog:

Email: paul@mrpaulandrew.com

Twitter: @mrpaulandrew

LinkedIn: In/mrpaulandrew

github.com/mrpaulandrew GitHub: